

The AI Generalist

Settings Lab: ChatGPT + Claude

Make the tools safer, more personal, and more useful before you prompt.

Mindset: We are not trying to become AI engineers today. We are becoming savvy users who know where the controls are, what they mean, and what not to upload.

Lab overview	Student-facing target
Time	60-90 minutes. Can be shortened by skipping the optional Claude Artifact extension.
Tools	ChatGPT and Claude only. Free accounts are enough for the core lab. Paid features are optional and should not be required.
Main outcome	Students leave with one strong prompt, one context packet, configured privacy/personalization settings, and one reusable Project/workspace in each tool.
Privacy posture	Use fake or harmless data. Do not upload student submissions, grades, accommodations, school records, private emails, client data, health data, financial data, or API keys.

Part 1 — User Prompt Engineering: Weak Prompt vs. Strong Prompt

Goal

You can see why prompt structure matters. This is the front door of AI Generalist work: clear delegation.

1. Open ChatGPT in one browser tab and Claude in a second browser tab.
2. Paste the weak prompt below into both tools. Do not improve it yet. Copy or screenshot the first response.

COPY / PASTE — Weak user prompt — Paste into the normal chat box

Help me make a lesson about AI.

3. Paste the strong prompt below into both tools in the same chat. Copy or screenshot the second response.

COPY / PASTE — Strong user prompt — Paste into the normal chat box

Act as an instructional designer helping a college professor build a 20-minute beginner lesson on “The AI Generalist.”

Audience: undergraduate or graduate students with mixed technical backgrounds.

Task: Create a lesson outline with 4 sections, one classroom analogy, one quick discussion question, and one 5-minute activity.

Constraints: keep language plain, avoid hype, and include one caution about privacy.

Output format: markdown table with columns: Section, Time, Teaching Point, Activity, Instructor Note.

4. Compare the outputs using the evaluation table. The goal is not to crown a winner. The goal is to see how better instructions change the answer.

Evaluation lens	What to look for	ChatGPT notes	Claude notes
Specificity	Does the model know the audience, time, task, and boundaries?		
Format control	Did it follow the requested table format?		
Teaching usefulness	Could a professor use the output with minimal editing?		
Privacy awareness	Did it include a useful caution, not generic filler?		
Style	Which output sounded most useful for teaching?		

Aha moments to ponder

Notice that the LLM did not suddenly become “smarter.” The assignment became less ambiguous. A strong prompt is like handing a colleague a well-written work order rather than a sticky note that says “do AI lesson.”

- Prompt engineering is not about magic words. It is about reducing guesswork.
- A strong prompt usually includes role, audience, task, constraints, examples or success criteria, and output format.
- If you change the prompt, you cannot fairly compare models. Same task first, then compare behavior.

Part 2 — Context Engineering: Give the Model the “Whole Folder,” Not Just the Question

Goal

You learn the difference between a better instruction and a better situation. Prompt engineering improves the order. Context engineering gives the assistant the background a capable human would need.

1. Paste the context packet below into ChatGPT and ask for the requested handout.
2. Paste the same context packet into Claude.
3. Now remove the context packet and ask only: “Create a handout about how to think like an AI Generalist.”
4. Compare what changed. Identify three details that were only followed well when context was provided.

COPY / PASTE — Context packet — Paste into the normal chat box

CONTEXT PACKET

Course: The AI Generalist

Audience: students learning to use ChatGPT and Claude beyond basic chatbot use.

Student anxiety: “I do not know which LLM to use, and I worry I am prompting wrong.”

Instructor goal: teach practical habits, not tool worship.

Tone: professor-like, clear, warm, and concrete.

Non-negotiables: include privacy cautions, free-tier alternatives, and one hands-on activity.

TASK

Using only the context above, create a 1-page student handout titled “How to Think Like an AI Generalist.”

Include: 5 habits, one model-selection rule, and one warning about sensitive data.

COPY / PASTE — Context self-audit prompt — Use before a high-value task

Before answering, inspect my context. Tell me:

1. What is clear enough to proceed?
2. What important information is missing?
3. What assumptions would you otherwise make?
4. What one question would most improve the output?


Then produce the requested handout.

Context ingredient	Question it answers	Example
Audience	Who will read or use the output?	“beginner undergraduate students”
Purpose	Why are we asking?	“prepare a 20-minute lesson”
Source material	What should the model rely on?	syllabus excerpt, rubric paragraph, transcript, notes
Examples	What does good look like?	sample table, tone sample, prior output
Constraints	What must be avoided?	no private data, no invented citations, no hype

- A long prompt is not automatically good context. Irrelevant context is noise.
- Context is like onboarding a new employee: audience, purpose, sources, examples, and boundaries come before the task.
- Good context should be reusable. If you keep pasting the same background, it probably belongs in a Project or Custom Instruction.


Part 3 — The Backend Layer: Settings Are Part of Prompt Engineering**Why we go into Settings now**

Beginners think AI behavior starts in the chat box. AI Generalists know that behavior also comes from settings: style, memory, data controls, connected apps, project instructions, and privacy choices. This is where system-style prompt engineering begins.



ChatGPT Settings screenshot: Personalization contains style/tone choices and Custom Instructions.

Paste durable instructions here, not in the ordinary message box.



Claude Settings screenshot: General/Profile includes “Instructions for Claude.” This is a durable behavior layer across chats, separate from a one-time user prompt.

Instruction layer	Where it lives	What it controls	Example
User prompt	Normal chat box	One task or one conversation	“Create a 20-minute lesson outline...”
Custom Instructions / Instructions for Claude	Settings / Profile / Personalization area	Durable behavior across many chats	“Use clear examples, warn me about privacy risks...”
Project Instructions	Inside a ChatGPT Project or Claude Project	Durable behavior for one workflow or course folder	“For this course project, coach prompt improvement and cite source files...”

Aha moment


Do not paste your system-style instructions into the chat box every time. That works once, but it does not configure the assistant. Put durable rules in Custom Instructions, Claude Instructions, or Project Instructions.

Part 4 — ChatGPT Settings Walkthrough


Goal

Inspect most of the ChatGPT Settings menu and make safe, class-friendly choices. If a setting is missing, write “not available in my account.”

1. Go to <https://chatgpt.com> and sign in using the account you are allowed to use for class.
2. Click your profile name or icon. Depending on the interface, it may be in the lower-left or upper-right area.
3. Click Settings.
4. Move through the menu from top to bottom. Do not rush. The settings panel is the cockpit.
5. Do not upgrade, connect private apps, upload private data, or change school-managed settings unless the instructor tells you to do so.



ChatGPT Data Controls screenshot: Locate model-improvement controls, shared links, archived chats, export, and deletion areas. For this lab, do not train on student or school data.



ChatGPT Apps screenshot: Connected apps can be powerful, but they expand the access surface. In this lab, inspect apps but do not connect new ones.

ChatGPT setting area	Numbered student action	Best-practice note
General / Appearance	Choose readable theme and display preferences.	Any accessible theme is fine. Comfort and readability beat aesthetics.
Notifications	Review email, desktop, and push notifications if visible.	Choose a quiet setup for class and screen sharing.
Personalization	Find style/tone options, Custom Instructions, Memory, and related preferences.	Use personalization for harmless preferences; be cautious with memory.
Custom Instructions	Paste the provided AI Generalist instruction in the Custom Instructions area.	This is system-style behavior. It does not belong in the normal prompt box.
Memory / Manage memories	Inspect what ChatGPT remembers if available. Delete anything sensitive or wrong.	For class data, avoid memory. Use Projects or Temporary Chat for bounded work.
Data Controls	Find “Improve the model for everyone” or equivalent wording and turn it OFF for class work.	Do not train on student data, school records, grades, or private communications.
Temporary Chat	Locate how to start a Temporary Chat if available.	Use for exploratory or sensitive one-off work that should not persist.
Apps / Connectors / Plugins	Review connected apps. Do not connect Gmail, Drive, Calendar, Slack, finances, LMS, or work systems during this lab.	Connected tools can expose live private data. Inspection is enough today.

ChatGPT setting area	Numbered student action	Best-practice note
Voice	Review voice and audio settings if visible.	Use headphones. Record classmates only with consent.
Billing / Usage / Plan	Observe your plan and limits.	Do not upgrade during class. Free tier is valid.
Storage / Library	Review what files or chat items may be stored.	Know where your uploads live and how to remove them later.
Safety / Security and login / Account	Check account security, active sessions, password/MFA options, and account email.	A powerful AI account may contain chats, files, and connected apps. Protect it.

COPY / PASTE — ChatGPT Custom Instructions — Paste in Settings > Personalization > Custom Instructions

You are my AI Generalist learning coach for this course.

How to respond to me:

- Use clear, beginner-friendly explanations.
- Give examples before abstract theory.
- When I ask for help, first identify the task, audience, context, format, and constraints.
- If my prompt is vague, diagnose what is missing before answering.
- Ask at most one clarifying question unless the task is truly ambiguous.
- Separate facts from assumptions and recommendations.
- Do not invent sources, citations, grades, policies, or data.
- Do not complete graded work for me; coach my process and help me improve my own work.
- Warn me if I am about to share private, student, client, employer, health, financial, or school data.
- Prefer concise structure: headings, lists, tables, examples, and a short next step.

My course goal:

I am learning to become an AI Generalist: someone who uses AI tools responsibly for prompting, context engineering, projects, research, writing, analysis, and workflow design.

1. After pasting the Custom Instructions, save them.
2. Start a new normal chat.
3. Ask: “What instructions are you following from my settings?”
4. If the answer does not reflect your instructions, go back and check whether you pasted them in the correct Settings field rather than the chat box.

Part 5 — Claude Settings Walkthrough

Goal

Configure Claude deliberately: profile instructions, privacy controls, capabilities, connectors, and project behavior. Claude's exact menus may vary by plan and organization.

1. Go to <https://claude.ai> and sign in using your class-approved account.
2. Click your initials, profile area, or account menu.
3. Click Settings.
4. In the left menu, move through the visible sections: General, Account, Privacy, Billing, Capabilities, Claude Code, Skills, Connectors, and Plugins if available.
5. Do not connect Google Drive, Gmail, Slack, GitHub, calendar tools, or school/work systems during this lab.

Claude setting area	Numbered student action	Best-practice note
General / Profile	Check display name and "What should Claude call you?" Then find "Instructions for Claude."	Paste system-style instructions here for durable behavior across chats.
Appearance / Preferences	Choose a comfortable theme and interface preferences.	Readable tools reduce friction and mistakes.
Account	Check which account and plan you are using.	Note which account you are using: personal, school, or work identity.
Privacy	Look for data-use or model-improvement wording such as "Help improve Claude." Turn it OFF for class work if visible.	Student or school data should not be used for model improvement. If admin-managed, record that.
Billing / Usage	Observe limits and plan. Do not upgrade in class.	Free tier is acceptable. Paid features are extensions, not requirements.
Capabilities	Check whether Artifacts, analysis tools, memory, web/search, or other features are visible.	Turn on Artifacts only if needed for the optional extension. Keep experimental tools cautious.
Claude Code	Observe whether it exists. Do not configure it unless the course is explicitly using coding tools.	This lab is AI Generalist setup, not software-agent setup.
Skills	Review available skills if visible. Do not install or enable unknown skills during class.	Skills can change Claude behavior. Treat them like add-ons with permissions.
Connectors	Review connected apps. Leave them disconnected unless using instructor-approved demo data.	Connectors can expose live external data.
Plugins	Review plugins if visible. Do not add new plugins during this setup lab.	A good operator inspects capabilities before using them.
Projects	Create a focused Project for this lab and paste project instructions.	Projects are the right place for workflow-specific instructions and knowledge.

COPY / PASTE — Claude Instructions — Paste in Claude Settings > General/Profile > Instructions for Claude

I am learning to become an AI Generalist. Help me use AI tools responsibly for prompting, context engineering, projects, research, writing, analysis, and workflow design.

How you should respond:

- Use plain language, concrete examples, and warm professor-like guidance.
- If my request is vague, identify missing context before answering.
- Coach my process rather than doing graded work for me.
- Separate source-backed facts, assumptions, and recommendations.
- Warn me before I share private, student, client, employer, school, health, financial, or sensitive data.
- Prefer clear structure: headings, numbered steps for procedures, and short tables when useful.
- End with one practical next step when appropriate.

1. After pasting the Claude Instructions, save them.
2. Start a new normal Claude chat.
3. Ask: “What instructions are you following from my settings?”
4. If Claude does not reflect the instructions, verify that the text was placed in Settings, not the normal message field.

Part 6 — Projects: Turning One-Off Chats into Workspaces

Why Projects matter

A normal chat is like a sticky note. A Project is like a course folder. It can hold files, chats, instructions, and sometimes knowledge. It gives the assistant a stable context boundary.

6A — ChatGPT Project setup

1. In the ChatGPT sidebar, click Projects or New Project if available.
2. Name the project: AI Generalist Prompt Lab.
3. Open the project menu, often shown as three dots near the project name.
4. Choose Project Settings, Instructions, or a similar instruction area.
5. Paste the project instruction below. This is not the same place as the normal chat prompt.
6. Add one harmless context source: a fake assignment prompt, a sample rubric paragraph, or the sample context block below.

COPY / PASTE — ChatGPT Project Instruction — Paste inside Project Settings / Instructions

You are the assistant for my AI Generalist Prompt Lab project.

Inside this project:

- Coach prompt engineering and context engineering.
- When improving a prompt, show the weak version, the improved version, and the reason the improved version is better.
- Use this structure whenever possible: Identity, Context, Task, Format, Constraints, Verification.
- Do not complete graded student submissions. Help the student plan, revise, test, and reflect.
- If source material is provided, clearly separate “from the source” from “my inference.”
- If data seems private or risky, stop and ask the student to remove or anonymize it.
- Keep outputs classroom-friendly, warm, and practical.

6B — Claude Project setup

1. In Claude, click Projects in the left sidebar if available.
2. Create a new Project named AI Generalist Prompt Lab.
3. Find Project Instructions, Custom Instructions, or the project knowledge/instructions area.
4. Paste the Claude Project instruction below.
5. Add the same harmless context source used in ChatGPT.
6. Run the same test prompt in both Projects and compare behavior.

COPY / PASTE — Claude Project Instruction — Paste inside Claude Project Instructions

You are my Claude assistant for the AI Generalist Prompt Lab.

Your role:

- Coach me to write better prompts and better context packs.
- Prefer careful reasoning, plain language, and practical examples.
- When I provide source material, use it directly and say when something is an inference.
- If the request involves a student assignment, help me understand and improve my process rather than writing the final submission for me.
- If I include private or sensitive information, tell me to remove or anonymize it.
- When improving a prompt, use this pattern: Identity, Context, Task, Format, Constraints, Verification.
- End with one concrete next step.

COPY / PASTE — Safe Sample Context Source — Add to both Projects

Sample assignment context for practice only:

Course: Introduction to AI Generalist Skills

Student level: beginner undergraduate students

Assignment: Students must create a one-page AI-assisted study guide from lecture notes.

Audience: classmates who are new to AI

Instructor priorities: clarity, ethical AI use, no private data, evidence-based claims, and concise formatting

Forbidden: do not invent citations, do not write the entire final submission for the student, and do not use personal student data.

COPY / PASTE — Project Test Prompt — Run inside both Projects

Use the project context to improve this weak prompt:

“Make me a study guide for AI.”

Return:

1. What is missing from the weak prompt
2. A stronger prompt using Identity, Context, Task, Format, Constraints, and Verification
3. A short explanation of why the new prompt is better
4. One privacy warning a student should remember

Part 7 — Optional Extension: Claude Artifacts as a Reusable Output

Use this only if time allows

Artifacts turn a chat answer into a working object: a checklist, mini-app, dashboard, quiz, calculator, diagram, or document. The prior Artifacts workbook emphasized a simple loop: describe what you want, see it run in a side panel, refine by chatting, and share when ready.

Gen AI Bootcamp Artifacts Workbook

3 • Build Your First Artifact


🕒 **The 15-minute lab.** Follow these six steps in order. By the end you will have a real finance-tracking app, live at a link you can send to anyone.

1. Go to Claude.ai

Open your browser and go to **claude.ai**. Sign in — a free account is enough for everything in this lab.

2. Open Settings → Capabilities

Click your profile in the bottom-left, choose **Settings**, then open the **Capabilities** tab in the left menu. Scroll down to the **Visuals** section near the bottom.



Settings → Capabilities — the three Visuals toggles you need are at the bottom.

3. Turn on the three Visuals toggles

Switch all three of these on:

- **Artifacts** — lets Claude build code, documents, and designs in a dedicated window.

Claude Artifacts workbook screenshot: Open Settings > Capabilities and enable Artifacts or related Visuals toggles before building an artifact.

1. Open Claude Settings.
2. Click Capabilities if the menu is visible.
3. Look for Artifacts, AI-powered artifacts, or Inline visualizations. Turn on only what the instructor approves for class.
4. Close Settings and start a new chat inside your AI Generalist Prompt Lab project.
5. Paste the artifact prompt below.
6. When the artifact appears, test it once and ask Claude for one improvement.

COPY / PASTE — Claude Artifact Prompt — Prompt Quality Checklist

Build an artifact called “Prompt Quality Checklist.”

Purpose: help students evaluate a user prompt before submitting it to ChatGPT or Claude.

Include five sections:

1. Identity: who should the assistant act as?
2. Context: what background does the assistant need?
3. Task: what exactly should happen?
4. Format: what should the output look like?
5. Constraints: what should be included, avoided, or verified?

Add a 1-5 score for each section, a “missing context” note box, and a final “next revision” box.

Design: clean, classroom-friendly, easy to read, and beginner-friendly.

Important: do not require login, database, private data, or external APIs.

- The artifact is not the main goal. The main goal is learning to describe a useful tool clearly.
- If Artifacts are unavailable, ask Claude for the same checklist as a plain table or markdown document.
- A good AI Generalist can downgrade gracefully when a feature is missing.

Part 8 — Compare Settings Behavior, Not Just Answers

Professional habit

The better question is not “Which model is best?” The better question is: “After I configured the settings correctly, which assistant behaved more safely and more usefully for this workflow?”

Criterion	ChatGPT evidence	Claude evidence	Winner / note
Did it follow custom/project instructions?			
Did it identify missing context?			
Did it avoid completing graded work for the student?			
Did it separate source facts from assumptions?			
Did it warn about privacy or sensitive data?			
Did the settings make the answer easier to use?			

Part 9 — Best Practices and What Not To Do

Do	Why	Do not
Use strong user prompts for one task.	They reduce ambiguity in the immediate request.	Do not type vague keywords and blame the model for guessing.
Use Custom Instructions / Claude Instructions for durable behavior.	They reduce repeated setup and make answers more consistent.	Do not paste system-style rules into every normal user prompt forever.
Use Projects for recurring workflows.	Projects keep context, files, chats, and instructions together.	Do not scatter the same course or client work across random chats.
Turn off training/model-improvement toggles for class data.	Student and school material should not become model-improvement material.	Do not assume consumer accounts are automatically configured for educational privacy.
Review memory before relying on it.	Memory can personalize responses but can also reuse unwanted context.	Do not let AI remember sensitive facts just because it can.
Use fake or anonymized data in class.	You can learn the workflow without exposing real people.	Do not upload grades, accommodations, advising notes, health info, private emails, or API keys.
Inspect connected apps before using them.	Connections can expose live files, emails, calendars, or financial data.	Do not connect Gmail, Drive, Slack, Calendar, LMS, or employer systems casually.
Save prompts that work.	Prompts become reusable assets.	Do not treat every good prompt as disposable.

- Accuracy: ask for sources or source references when facts matter, then spot-check the answer.
- Privacy: settings can reduce risk, but they do not make unsafe data magically safe.
- Accountability: AI can draft and organize. Humans remain responsible for what is sent, submitted, graded, published, or deployed.

Part 10 — Final Submission: AI Generalist Setup Ticket

1. Submit one weak prompt and the strong prompt you wrote.
2. Submit a screenshot or written note showing where you pasted Custom Instructions in ChatGPT.
3. Submit a screenshot or written note showing where you pasted Instructions for Claude.
4. Submit one sentence explaining which model followed your settings better and why.
5. Submit one privacy rule you will follow for the rest of the semester.

Setup ticket question	Student response
ChatGPT: Did you find the model-improvement/data-use toggle?	
ChatGPT: Did you turn it off for class data or record why it was unavailable/admin-managed?	
ChatGPT: Where did you paste your system-style instruction?	
ChatGPT: Did you create a Project? What did you name it?	
Claude: Did you find the privacy/data improvement setting?	
Claude: Where did you paste your durable instruction?	
Claude: Did you create a Project? What did you name it?	
What is one setting you did not know existed before today?	
What is one thing you will never paste into an AI chat?	

Mini-rubric

Criterion	Pts	Full-credit evidence
Strong prompt comparison	2	Student compares weak vs strong prompt in both ChatGPT and Claude.
Context engineering	2	Student provides a useful context packet and explains how output changed.
Settings literacy	3	Student inspects privacy, personalization, memory, data controls, and connected apps in both tools.
System-style instructions	2	Student pastes durable instructions in the correct Settings/Project field, not the normal chat box.
Reflection	1	Student names one risk, one useful setting, and one behavior they will keep.